WILLIAMS-SONOMA, INC.

AI-Thon 2026

Free AI Tools Guide

Everything you need to build. No credit card required.

A few things before you start

- You are welcome to use paid tools or upgrade to a paid plan if you want to. Just know that Williams-Sonoma will not reimburse those costs. That is entirely your call and your expense. Free tiers are what we recommend, and everything in this guide is based on what you can get for free.
- If a tool starts throttling you, add a short retry or switch to one of the backups listed here. It happens.
- Write down every AI tool you used in your README. Judges will ask, and it helps them understand your work.
- Keep API keys out of your repo. Use a .env file and add it to .gitignore from the start.

1. AI Chat and Ideation

These are your go-to tools for thinking through ideas, writing prompts, debugging logic, and tidying up your pitch. Keep one open in a tab all weekend.

Tool	Category	Good for	Free tier limits
ChatGPT (OpenAI)	AI chat	Ideation, code help, writing, debugging	Free: GPT-4o mini; GPT-4o has a daily message limit
Claude (Anthropic)	AI chat	Long documents, prompt design, writing	Free at claude.ai; no free API access
Gemini (Google)	AI chat	Multimodal work; connects well with Google tools	Free at gemini.google.com with solid limits
Perplexity AI	AI search	Research with citations; good for domain deep dives	Free with unlimited standard searches

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.

Perplexity is really useful for sector research like healthcare regulations or fintech compliance basics. Use ChatGPT or Claude for the actual coding and writing.

2. LLM APIs: Adding AI to Your App

This is how you give your app a brain. Call these APIs from your backend to power chat, summarization, document Q&A, agents, and more.

Tool	Category	Good for	Free tier limits
Google AI Studio	LLM API	Gemini 1.5 Flash and Pro; handles text, images, audio, video	15 requests/min, 1 million tokens/day free
Gemini CLI	Terminal AI	Run Gemini from your terminal; handy for scripting	Free with a Google account; shares AI Studio quota
OpenAI Codex (API)	Code LLM	Code generation, explanation, and completion	Free credits on new accounts; worth checking current offer
Groq	LLM API	Very fast Llama 3 and Mixtral responses	14,400 requests/day free; best option for low latency
Mistral (La Plateforme)	LLM API	Mistral 7B and Mixtral; solid at reasoning tasks	Free credits when you sign up
OpenRouter (free models)	LLM routing	One API key to reach 50+ models	Free for any model tagged :free

A sensible default: Google AI Studio for anything needing lots of tokens, Groq when speed matters, and OpenRouter as a fallback if you want to try a different model quickly.

3. AI Coding Assistants and IDEs

These plug into your editor and help you write, fix, and refactor code faster. Cursor is where most teams land on the free tier.

Tool	Category	Good for	Free tier limits
Cursor	AI-first IDE	Edits across multiple files; inline chat; ask about your whole codebase	Free tier: 2,000 completions/month
VS Code + Copilot	IDE extension	Inline suggestions, chat panel, terminal agent	Free for students through GitHub Education
GitHub Copilot	Code completion	Autocomplete in VS Code, JetBrains, and Neovim	Free for verified students

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.

OpenAI Codex (CLI)	Terminal agent	Describe what you want in plain English; it edits your files	Free credits on new accounts
Aider	Terminal agent	AI pair programming in the terminal; works across files	Open source; pair with Groq or Gemini for free access
Replit	Browser IDE	Code and run in a browser tab; good if local setup is tricky	Free tier with usage limits

4. Full-Stack and UI Builders

Describe your app in plain English and get working React or Next.js code back. A great way to get moving fast in the first hour, especially if your team is light on frontend experience.

Tool	Category	Good for	Free tier limits
Bolt.new	Full-stack builder	Builds React and Node apps from a description; you can keep editing	Around 150k tokens/day free
Lovable.dev	App builder	UI-first React apps with GitHub export built in	Free tier with a project limit per month
Vo.dev (Vercel)	UI component gen	Generates React components using shadcn/ui	200 credits/month free
Antigravity	AI app builder	No-code and code-friendly; build and deploy from one place	Free tier available; check current limits on their site

One approach that works well: use Bolt.new or Antigravity to get a working scaffold in the first hour, then take it into Cursor for the rest of the build. Saves a lot of boilerplate time.

5. Version Control and Collaboration (Required)

Every team needs a GitHub repo. It is part of your submission and judges will look at your commit history, so start it early.

Tool	Category	Good for	Free tier limits
GitHub	Code repo	Your submission repo; handles CI/CD through Actions	Free for public and private repos with no limits
GitHub Codespaces	Cloud IDE	VS Code in the browser; useful if you cannot install things locally	Free: 60 hours/month per account
GitHub Actions	CI/CD	Auto-deploy on push; run tests; basic build pipelines	Free: 2,000 minutes/month

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.

Create your team repo in the first 30 minutes. All code goes there. Paste the link in your submission form when you are done.

6. Sharing Your App and Local Tunnelling

ngrok lets you share an app running on your laptop with anyone in the world via a public URL. Really useful during the build phase and a lifesaver for last-minute demo prep.

Tool	Category	Good for	Free tier limits
ngrok	Local tunnel	Share localhost as a live HTTPS URL with one command	Free: 1 tunnel, no credit card needed
Vercel	Frontend host	Deploy React and Next.js apps; auto-deploys from GitHub	Free on the hobby plan with unlimited deploys
Railway	Backend host	Deploy APIs, databases, and background workers	$5 in free credits per month
Supabase	Database + Auth	Postgres database, auth, file storage, and edge functions	Free: 500MB database, 1GB storage
Firebase	Backend + host	Realtime database, auth, storage, and hosting	Free Spark plan with generous limits
Hugging Face Spaces	App hosting	Host Gradio or Streamlit demos; deploy ML models	Free on CPU instances

Getting ngrok running: sign up at ngrok.com (free), run npm install -g ngrok, then ngrok http 3000. You get a live public URL in about 30 seconds.

7. Specialist AI Tools

Tool	Category	Good for	Free tier limits
Whisper (OpenAI)	Speech to text	Transcribes audio and video; open source	Free; run locally or use a free community API wrapper
ElevenLabs	Text to speech	Good voice synthesis; works well in demos	10,000 characters/month free
Stability AI	Image generation	Text to image via Stable Diffusion	Free credits when you sign up
Hugging Face Models	Model hub	NLP, vision, and audio models through the Inference API	Free tier with rate limits
LangChain	AI framework	Build RAG pipelines, agents, and chains in any language	Open source and free

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.

LlamaIndex	AI framework	Document Q&A and retrieval-augmented generation	Open source and free

8. Languages and Recommended Stacks

You can build in whatever language you are comfortable with. The ones below are preferred because most of our mentors know them well and the tools above integrate with them most easily.

Language / Stack	Why it works well at AI-Thon
React.js (preferred)	Component-based frontend that works seamlessly with Vercel, Bolt.new, V0.dev, and Lovable
Next.js (preferred)	Combines React with server-side logic and API routes in a single project; one-click deploy on Vercel
Node.js (preferred)	JavaScript on the backend; quick to wire up REST APIs and connect to LLM APIs
Python (preferred)	The natural home for AI-heavy work: LangChain, LlamaIndex, Hugging Face, data tasks, and Streamlit demos
JavaScript	Works great for lightweight frontends or browser extensions where you do not need a build step
Java	Solid for enterprise integrations and Android backends; Spring Boot deploys well on Railway
Flutter	Cross-platform mobile and web; a good fit for Healthcare or Social Impact projects with a mobile focus
Any other language	Totally fine. Just make sure your README explains how to run the project so judges can try it themselves

If you are not sure where to start: Next.js for the frontend and API layer, Python for AI logic, Supabase for the database, and Vercel for hosting. Our mentors know this stack well.

Useful Tips

Staying under rate limits

- Cache responses during development so you are not burning API calls while tweaking your UI
- Groq is fast; Google AI Studio handles volume. Use them for what they are each good at.
- A simple retry with a short wait handles most rate limit errors without any drama
- Mock your AI responses while building the UI. Switch to real calls when you are ready to test end to end.

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.

Sharing with teammates and judges

- Get ngrok running from Hour 1. It stops the "works on my machine" problem before it starts.
- Aim to have something deployed on Vercel by Hour 30. The last stretch should be for polish, not fixing deployments.
- Make sure a judge can reach your core feature in under 30 seconds with no instructions from you.

Good habits with GitHub

- Commit regularly, at least every couple of hours. Judges do look at commit history.
- Your README matters. Cover the problem you solved, a demo link, the tech stack, how to run it, and what AI tools you used.
- Use a .env.example file to document your environment variables without putting real keys in the repo.

© 2026 Williams-Sonoma, Inc. All rights reserved. For internal and event use only.